


Numerical Integration

Computer Graphics
CMU 15-462/15-662

Motivation: The Rendering Equation

- Recall the rendering equation, which models light “bouncing around the scene”:



$$L_o(\mathbf{p}, \omega_o) = L_e(\mathbf{p}, \omega_o) + \int_{\mathcal{H}^2} f_r(\mathbf{p}, \omega_i \rightarrow \omega_o) L_i(\mathbf{p}, \omega_i) \cos \theta d\omega_i$$


How can we possibly evaluate this integral?

Numerical Integration—Overview

- In graphics, many quantities we're interested in are naturally expressed as integrals (total brightness, total area, ...)
- For very, very simple integrals, we can compute the solution analytically
- For everything else, we have to compute a numerical approximation
- Basic idea:
 - integral is “area under curve”
 - sample the function at many points
 - integral is approximated as weighted sum



$$\int_0^1 \frac{1}{3}x^2 \, dx = \left[x^3 \right]_0^1 = 1$$




Rendering: what are we integrating?

- Recall this view of the world:




Want to “sum up”—i.e., integrate!—light from all directions
(But let’s start a little simpler...)

Review: integral as “area under curve”



Or: average value times size of domain


$$\int_a^b f(x)dx = (b - a)\text{mean}(f)$$



Review: fundamental theorem of calculus


$$\int_a^b f(x)dx = F(b) - F(a)$$

$$f(x) = \frac{d}{dx} F(x)$$




Simple case: constant function

$$\int_a^b C dx = (b - a)C$$




Affine function: $f(x) = cx + d$

$$\int_a^b f(x)dx = \frac{1}{2}(f(a) + f(b))(b - a)$$




Need only one sample of the function (at just the right place...)

More general polynomials?




Gauss Quadrature


- For any polynomial of degree $2n-1$ or less, we can always obtain the exact integral by sampling at a special set of n points and taking a special weighted combination




$n=1$




$n=2$



$n=4$




$n=3$



$n=5$

Piecewise affine function

For piecewise functions, just sum integral of each piece:




Key idea so far:

To approximate an integral, we need

- (i) quadrature points, and**
- (ii) weights for each point**

$$\int_a^b f(x) dx \approx \sum_{i=1}^n w_i f(x_i)$$

Arbitrary function $f(x)$?




Trapezoid rule

Approximate integral of $f(x)$ by pretending function is piecewise affine

For equal length segments:

$$h = \frac{b - a}{n - 1}$$

$$\int_a^b f(x) dx = h \left(\sum_{i=1}^{n-1} f(x_i) + \frac{1}{2} (f(x_0) + f(x_n)) \right)$$




Trapezoid rule

Consider cost and accuracy of estimate as $n \rightarrow \infty$ (or $h \rightarrow 0$)


Work: $O(n)$

Error can be shown to be: $O(h^2) = O(\frac{1}{n^2})$

(for $f(x)$ with continuous second derivative)



What about a 2D function?



How should we approximate the area underneath?

Integration in 2D

Consider integrating $f(x, y)$ using the trapezoidal rule
(apply rule twice: when integrating in x and in y)

$$\begin{aligned} \int_{a_y}^{b_y} \int_{a_x}^{b_x} f(x, y) dx dy &= \int_{a_y}^{b_y} \left(O(h^2) + \sum_{i=0}^n A_i f(x_i, y) \right) dy \quad \text{First application of rule} \\ &= O(h^2) + \sum_{i=0}^n A_i \int_{a_y}^{b_y} f(x_i, y) dy \\ &= O(h^2) + \sum_{i=0}^n A_i \left(O(h^2) + \sum_{j=0}^n A_j f(x_i, y_j) \right) \quad \text{Second application} \\ &= O(h^2) + \sum_{i=0}^n \sum_{j=0}^n A_i A_j f(x_i, y_j) \end{aligned}$$

Errors add, so error still: $O(h^2)$

But work is now: $O(n^2)$
($n \times n$ set of measurements)

Must perform much more work in 2D to get same error bound on integral!


In K-D, let $N = n^k$

Error goes as: $O\left(\frac{1}{N^{2/k}}\right)$


Curse of Dimensionality

- How much does it cost to apply the trapezoid rule as we go up in dimension?

- 1D: $O(n)$
- 2D: $O(n^2)$
- ...
- $kD: O(n^k)$



- For many problems in graphics (like rendering), k is very, very big (e.g., tens or hundreds or thousands)
- Applying trapezoid rule does not scale!
- Need a fundamentally different approach...



Monte Carlo Integration

Monte Carlo Integration

So far we've discussed techniques that use a fixed set of sample points (e.g., uniformly spaced, or obtained by finding roots of polynomial (Gaussian quadrature))

- Estimate value of integral using random sampling of function
 - Value of estimate depends on random samples used
 - But algorithm gives the correct value of integral “on average”
- Only requires function to be evaluated at random points on its domain
 - Applicable to functions with discontinuities, functions that are impossible to integrate directly
- Error of estimate is independent of the dimensionality of the integrand
 - Depends on the number of random samples used: $O\left(\frac{1}{\sqrt{n}}\right)$

Review: random variables

X

random variable. Represents a distribution of potential values

$X \sim p(x)$ **probability density function (PDF). Describes relative probability of a random process choosing value x**

Uniform PDF: all values over a domain are equally likely

e.g., for an unbiased die

X takes on values 1,2,3,4,5,6

$p(1) = p(2) = p(3) = p(4) = p(5) = p(6)$



Discrete probability distributions

n discrete values x_i

With probability p_i


Requirements of a PDF:

$$p_i \geq 0$$

$$\sum_{i=1}^n p_i = 1$$

Six-sided die example: $p_i = \frac{1}{6}$

Think: p_i is the probability that a random measurement of X will yield the value x_i
 X takes on the value x_i with probability p_i



Cumulative distribution function (CDF)



(For a discrete probability distribution)

Cumulative PDF: $P_j = \sum_{i=1}^j p_i$

where:

$$0 \leq P_i \leq 1$$

$$P_n = 1$$



How do we generate samples of a discrete random variable (with a known PDF?)



Sampling from discrete probability distributions

To randomly select an event, select x_i if

$$P_{i-1} < \xi \leq P_i$$



Uniform random variable $\in [0, 1)$



Continuous probability distributions

PDF $p(x)$

$$p(x) \geq 0$$

CDF $P(x)$

$$P(x) = \int_0^x p(x) \, dx$$



$$P(x) = \Pr(X < x)$$

$$P(1) = 1$$

$$\begin{aligned} \Pr(a \leq X \leq b) &= \int_a^b p(x) \, dx \\ &= P(b) - P(a) \end{aligned}$$

Uniform distribution

(for random variable X defined on $[0,1]$ domain)



Sampling continuous random variables using the inversion method

Cumulative probability distribution function


$$P(x) = \Pr(X < x)$$

Construction of samples:

Solve for $x = P^{-1}(\xi)$


Must know the formula for:

1. The integral of $p(x)$
2. The inverse function $P^{-1}(\xi)$




Example—Sampling Quadratic Distribution

- As a toy example, consider the simple probability distribution $p(x) := 3(1-x)^2$ over the interval $[0,1]$
- How do we pick random samples distributed according to $p(x)$?
- First, integrate probability distribution $p(x)$ to get cumulative distribution $P(x)$
- Invert $P(x)$ by solving $\xi = P(x)$ for x
- Finally, plug uniformly distributed random values ξ in $[0,1]$ into this expression




$$\int_0^x 3(1-s)^2 ds = x^3 - 3x^2 + 3x$$

$$x = P^{-1}(\xi) = 1 - (1 - \xi)^{\frac{1}{3}}$$



How do we uniformly sample the unit circle?



I.e., choose any point $P=(px, py)$ in circle with equal probability)


Uniformly sampling unit circle: first try

- $\theta = \text{uniform random angle between } 0 \text{ and } 2\pi$
- $r = \text{uniform random radius between } 0 \text{ and } 1$
- Return point: $(r \cos \theta, r \sin \theta)$

This algorithm does not produce the desired uniform sampling of the area of a circle. Why?

Because sampling is not uniform in area!

Points farther from center of circle are less likely to be chosen



$$\theta = 2\pi\xi_1 \quad r = \xi_2$$

So how should we pick samples? Well, think about how we integrate over a disk in polar coordinates...

Sampling a circle (via inversion in 2D)

$$A = \int_0^{2\pi} \int_0^1 r \, dr \, d\theta = \int_0^1 r \, dr \int_0^{2\pi} d\theta = \left(\frac{r^2}{2} \right) \Big|_0^1 \Big|_0^{2\pi} = \pi$$

$$p(r, \theta) \, dr \, d\theta = \frac{1}{\pi} r \, dr \, d\theta \rightarrow p(r, \theta) = \frac{r}{\pi}$$

$$p(r, \theta) = p(r)p(\theta) \quad \leftarrow \quad r, \theta \text{ independent}$$

$$p(\theta) = \frac{1}{2\pi}$$

$$P(\theta) = \frac{1}{2\pi}\theta$$

$$p(r) = 2r$$

$$P(r) = r^2$$


$$\xi_1 = P(\theta) = \frac{\theta}{2\pi}$$

$$\theta = 2\pi\xi_1$$

$$\xi_2 = P(r) = r^2$$

$$r = \sqrt{\xi_2}$$


so that we integrate to
1 instead of area



Uniform area sampling of a circle

WRONG

**probability is uniform;
samples are not!**




$$\theta = 2\pi\xi_1$$

$$r = \xi_2$$

RIGHT

**probability is nonuniform;
samples are uniform**




$$\theta = 2\pi\xi_1$$

$$r = \sqrt{\xi_2}$$

Uniform sampling via rejection sampling


**Completely different idea: pick uniform samples in square (easy)
Then toss out any samples not in square (easy)**



Efficiency of technique: area of circle / area of square

Efficiency of Rejection Sampling

- If the region we care about covers only a very small fraction of the region we're sampling, rejection is probably a bad idea:



Smarter in this case to “warp” our random variables to follow the spiral.

**So how do we use numerical integration to
do rendering?**

Monte Carlo Rendering

- Goal: render a photorealistic image
- Put together many of the ideas we've studied:
 - color
 - materials
 - radiometry
 - numerical integration
 - geometric queries
 - spatial data structures
 - rendering equation
- Combine into final Monte Carlo ray tracing algorithm
- Alternative to rasterization, lets us generate much more realistic images (usually at much greater cost...)




Photorealistic Rendering—Basic Goal

What are the **INPUTS** and **OUTPUTS**?

camera



geometry



materials



lights



Ray Tracer

("scene")




image

Ray Tracing vs. Rasterization—Order

- Both rasterization & ray tracing will generate an image
- What's the difference?
- One basic difference: order in which we process samples

RASTERIZATION



for each primitive:


for each sample:

 determine coverage

 evaluate color

**(Use Z-buffer to determine
which primitive is visible)**

RAY TRACING



for each sample:

for each primitive:

 determine coverage


 evaluate color

**(Use spatial data structure like BVH to
determine which primitive is visible)**


Ray Tracing vs. Rasterization—Illumination

- More major difference: sophistication of illumination model
 - [LOCAL] rasterizer processes one primitive at a time; hard* to determine things like “A is in the shadow of B”
 - [GLOBAL] ray tracer processes on ray at a time; ray knows about everything it intersects, easy to talk about shadows & other “global” illumination effects

RASTERIZATION



RAY TRACING




Q: What illumination effects are missing from the image on the left?

*But not impossible to do some things with rasterization (e.g., shadow maps)... just results in more complexity

Monte Carlo Ray Tracing

- To develop a full-blown photorealistic ray tracer, will need to apply Monte Carlo integration to the rendering equation
- To determine color of each pixel, integrate incoming light
- What function are we integrating?
 - illumination along different paths of light
- What does a “sample” mean in this context?
 - each path we trace is a sample

$$L_o(\mathbf{p}, \omega_o) = L_e(\mathbf{p}, \omega_o) + \int_{\mathcal{H}^2} f_r(\mathbf{p}, \omega_i \rightarrow \omega_o) L_i(\mathbf{p}, \omega_i) \cos \theta d\omega_i$$



Monte Carlo Integration

- Started looking at Monte Carlo integration in our lecture on numerical integration
- Basic idea: take average of random samples
- Will need to flesh this idea out with some key concepts:
 - EXPECTED VALUE — what value do we get on average?
 - VARIANCE — what's the expected deviation from the average?
 - IMPORTANCE SAMPLING — how do we (correctly) take more samples in more important regions?

$$\lim_{N \rightarrow \infty} \frac{|\Omega|}{N} \sum_{i=1}^N f(X_i) = \int_{\Omega} f(x) dx$$

Expected Value

Intuition: what value does a random variable take, on average?

- E.g., consider a fair coin where heads = 1, tails = 0
- Equal probability of heads & is tails (1/2 for both)
- Expected value is then $(1/2) \cdot 1 + (1/2) \cdot 0 = 1/2$

$$E(Y) := \sum_{i=1}^k p_i y_i$$

expected value of random variable Y

number of possible outcomes

probability of ith outcome

value of ith outcome

Properties of expectation:

$$E \left[\sum_i Y_i \right] = \sum_i E[Y_i]$$
$$E[aY] = aE[Y]$$

(Can you show these are true?)



Variance

Intuition: how far are our samples from the average, on average?

Definition

$$V[Y] = E[(Y - E[Y])^2]$$

Q: Which of these has higher variance?



Properties of variance:

$$V[Y] = E[Y^2] - E[Y]^2$$

$$V \left[\sum_{i=1}^N Y_i \right] = \sum_{i=1}^N V[Y_i]$$

$$V[aY] = a^2 V[Y]$$

(Can you show these are true?)


Law of Large Numbers

- Important fact: for any random variable, the average value of N trials approaches the expected value as we increase N
- Decrease in variance is always linear in N :

$$V \left[\frac{1}{N} \sum_{i=1}^N Y_i \right] = \frac{1}{N^2} \sum_{i=1}^N V[Y_i] = \frac{1}{N^2} N V[Y] = \frac{1}{N} V[Y]$$

Consider a coconut...

nCoconuts	estimate of π
1	4.000000
10	3.200000
100	3.240000
1000	3.112000
10000	3.163600
100000	3.139520
1000000	3.141764




Q: Why is the law of large numbers important for Monte Carlo ray tracing?

A: No matter how hard the integrals are (crazy lighting, geometry, materials, etc.), can always* get the right image by taking more samples.

***As long as we make sure to sample all possible kinds of light paths...**

Biassing


- So far, we've picked samples uniformly from the domain (every point is equally likely)
- Suppose we pick samples from some other distribution (more samples in one place than another)



- Q: Can we still use samples $f(X_i)$ to get a (correct) estimate of our integral?
- A: Sure! Just weight contribution of each sample by how likely we were to pick it
- Q: Are we correct to divide by p ? Or... should we multiply instead?

$$\int_{\Omega} f(x) dx \approx \frac{1}{N} \sum_{i=1}^N \frac{f(X_i)}{p(X_i)}$$

- A: Think about a simple example where we sample RED region 8x as often as BLUE region
 - average color over square should be purple
 - if we multiply, average will be TOO RED
 - if we divide, average will be JUST RIGHT



Next Time: Use biasing for Importance Sampling, along with other aspects of effective Monte Carlo Raytracing!

